

Package ‘fastrerandomize’

September 26, 2025

Title Hardware-Accelerated Rerandomization
for Improved Balance

Version 0.2

Description Provides hardware-accelerated tools for performing rerandomization and randomization testing in experimental research. Using a 'JAX' backend, the package enables exact rerandomization inference even for large experiments with hundreds of billions of possible randomizations. Key functionalities include generating pools of acceptable rerandomizations based on covariate balance, conducting exact randomization tests, and performing pre-analysis evaluations to determine optimal rerandomization acceptance thresholds. The package supports various hardware acceleration frameworks including 'CPU', 'CUDA', and 'METAL', making it versatile across accelerated computing environments. This allows researchers to efficiently implement stringent rerandomization designs and conduct valid inference even with large sample sizes. The package is partly based on Jerzak and Goldstein (2023) <doi:10.48550/arXiv.2310.00861>.

URL <https://github.com/cjerzak/fastrerandomize-software>

BugReports <https://github.com/cjerzak/fastrerandomize-software/issues>

Depends R (>= 3.3.3)

License GPL-3

Encoding UTF-8

LazyData false

Imports reticulate,
assertthat,
utils,
stats,
graphics

Suggests knitr,
rmarkdown

VignetteBuilder knitr

RoxygenNote 7.3.2

RemoteType local

RemotePkgRef local::~~/Documents/fastrerandomize-software/fastrerandomize

RemoteUrl /Users/cjerzak/Documents/fastrerandomize-software/fastrerandomize

Contents

build_backend	2
check_jax_availability	3
compute_diff_at_tau_for_oneW_R	4
diff_in_means_R	4
fastrandomize_class	5
fastrandomize_test	5
fast_distance	6
find_fiducial_interval_R	8
generate_randomizations	9
generate_randomizations_exact	11
generate_randomizations_mc	13
generate_randomizations_R	15
hotellingT2_R	16
plot.fastrandomize_randomizations	16
plot.fastrandomize_test	17
print.fastrandomize_randomizations	17
print.fastrandomize_test	18
print2	18
QJEData	19
randomization_test	20
randomization_test_R	22
summary.fastrandomize_randomizations	23
summary.fastrandomize_test	23
YOPData	24

Index	25
--------------	-----------

build_backend	<i>A function to build the environment for fastrandomize. Builds a conda environment in which 'JAX' and 'np' are installed. Users can also create a conda environment where 'JAX' and 'np' are installed themselves.</i>
---------------	--

Description

A function to build the environment for fastrandomize. Builds a conda environment in which 'JAX' and 'np' are installed. Users can also create a conda environment where 'JAX' and 'np' are installed themselves.

Usage

```
build_backend(conda_env = "fastrandomize_env", conda = "auto")
```

Arguments

conda_env	(default = "fastrandomize_env") Name of the conda environment in which to place the backends.
conda	(default = auto) The path to a conda executable. Using "auto" allows reticulate to attempt to automatically find an appropriate conda binary.

Value

Invisibly returns NULL; this function is used for its side effects of creating and configuring a conda environment for fastrerandomize. This function requires an Internet connection. You can find out a list of conda Python paths via: `Sys.which("python")`

Examples

```
## Not run:
# Create a conda environment named "fastrerandomize_env"
# and install the required Python packages (jax, numpy, etc.)
build_backend(conda_env = "fastrerandomize_env", conda = "auto")

# If you want to specify a particular conda path:
# build_backend(conda_env = "fastrerandomize_env", conda = "/usr/local/bin/conda")

## End(Not run)
```

check_jax_availability

Check if 'Python' and 'JAX' are available

Description

This function checks if 'Python' and 'JAX' can be accessed via 'reticulate'. If not, it returns 'NULL' and prints a message suggesting to run 'build_backend()'.

Usage

```
check_jax_availability(conda_env = "fastrerandomize_env", conda = "auto")
```

Arguments

conda_env	A character string specifying the name of the conda environment. Default is "fastrerandomize_env".
conda	The path to a conda executable, or "auto". Default is "auto".

Value

Returns 'TRUE' (invisibly) if both 'Python' and 'JAX' are available; otherwise returns 'NULL'.

Examples

```
## Not run:
  check_jax_availability()

## End(Not run)
```

```
compute_diff_at_tau_for_oneW_R
```

Compute potential outcome difference in means for a single assignment under a hypothesized tau in base R

Description

Compute potential outcome difference in means for a single assignment under a hypothesized tau in base R

Usage

```
compute_diff_at_tau_for_oneW_R(Wprime, obsY, obsW, tau)
```

Arguments

Wprime	A 0/1 assignment vector for which to compute the diff in means.
obsY	Observed outcome vector.
obsW	Observed assignment vector.
tau	The hypothesized true effect for the shift in outcomes under treatment.

Value

Scalar difference in means for the assignment Wprime.

```
diff_in_means_R
```

Simple difference in means in base R

Description

Simple difference in means in base R

Usage

```
diff_in_means_R(Y, W)
```

Arguments

Y	Numeric outcome vector.
W	0/1 treatment assignment vector.

Value

Scalar difference in means: $\text{mean}(Y|W=1) - \text{mean}(Y|W=0)$.

`fastrerandomize_class` *Constructor for fastrerandomize randomizations*

Description

Create an S3 object of class `fastrerandomize_randomizations` that stores the randomizations (and optionally balance statistics) generated by functions such as [generate_randomizations](#).

Usage

```
fastrerandomize_class(
  randomizations,
  balance = NULL,
  fastrr_env = NULL,
  call = NULL
)
```

Arguments

<code>randomizations</code>	A matrix or array where each row (or slice) represents one randomization.
<code>balance</code>	A numeric vector or similar object holding balance statistics for each randomization, or NULL if not applicable.
<code>fastrr_env</code>	Associated <code>fastrr_env</code> environment.
<code>call</code>	The function call, if you wish to store it for reference (optional).

Value

An object of class `fastrerandomize_randomizations`.

`fastrerandomize_test` *Constructor for fastrerandomize randomization test objects*

Description

Constructor for `fastrerandomize` randomization test objects

Usage

```
fastrerandomize_test(p_value, FI, tau_obs, fastrr_env = NULL, call = NULL, ...)
```

Arguments

<code>p_value</code>	A numeric value representing the p-value of the test.
<code>FI</code>	A numeric vector (length 2) representing the fiducial interval, or NULL if not requested.
<code>tau_obs</code>	A numeric value (or vector) representing the estimated treatment effect.
<code>fastrr_env</code>	Associated <code>'fastrr_env'</code> environment.
<code>call</code>	An optional function call, stored for reference.
<code>...</code>	Other slots you may want to store (e.g. additional diagnostics).

Value

An object of class `fastrerandomize_test`.

<code>fast_distance</code>	<i>JAX-accelerated distance calculations</i>
----------------------------	--

Description

Compute pairwise distances between the rows of one matrix A or two matrices A and B, using JAX-backed, JIT-compiled kernels. Supports common metrics: Euclidean, squared Euclidean, Manhattan, Chebyshev, Cosine, Minkowski (with optional feature weights), and Mahalanobis (with full or diagonal inverse covariance).

The function automatically batches computations to avoid excessive device memory use.

Usage

```
fast_distance(
  A,
  B = NULL,
  metric = c("euclidean", "sqeuclidean", "manhattan", "chebyshev", "cosine", "minkowski",
    "mahalanobis"),
  p = 2,
  weights = NULL,
  cov_inv = NULL,
  approximate_inv = TRUE,
  squared = FALSE,
  row_batch_size = NULL,
  as_dist = FALSE,
  return_type = "R",
  verbose = FALSE,
  conda_env = "fastrerandomize_env",
  conda_env_required = TRUE
)
```

Arguments

<code>A</code>	A numeric matrix with rows as observations and columns as features.
<code>B</code>	Optional numeric matrix with the same number of columns as A. If NULL, distances are computed within A (i.e., $n \times n$).
<code>metric</code>	Character; one of "euclidean", "sqeuclidean", "manhattan", "chebyshev", "cosine", "minkowski", "mahalanobis". Default is "euclidean".
<code>p</code>	Numeric order for Minkowski distance (must be > 0). Default is 2.
<code>weights</code>	Optional numeric vector of length <code>ncol(A)</code> with nonnegative feature weights. Used for "minkowski" and "manhattan" (the latter is equivalent to Minkowski with $p = 1$).
<code>cov_inv</code>	Optional inverse covariance matrix ($p \times p$) for Mahalanobis (ignored if <code>approximate_inv = TRUE</code>). If not supplied and <code>approximate_inv = FALSE</code> , it is estimated from <code>rbind(A, B)</code> and inverted in JAX.

approximate_inv	Logical; if TRUE and metric = "mahalanobis", uses a diagonal inverse (reciprocal variances) for speed and robustness. Default TRUE.
squared	Logical; if TRUE, return squared distances when supported ("euclidean" and "mahalanobis"). Ignored for other metrics. Default FALSE.
row_batch_size	Optional integer; number of rows of A to process per batch. If NULL, a safe size is chosen automatically.
as_dist	Logical; if TRUE and B is NULL, return a base dist object (for symmetric metrics). Default FALSE.
return_type	Either "R" (convert to base R matrix/dist) or "jax" (return a JAX array). Default "R".
verbose	Logical; print batching progress. Default FALSE.
conda_env	Character; conda environment name used by reticulate. Default "fastrerandomize_env".
conda_env_required	Logical; whether the specified conda environment must be used. Default TRUE.

Details

- **Mahalanobis**: with approximate_inv = TRUE, the diagonal of the pooled covariance is used (variance stabilizer); otherwise a full inverse covariance is used. - **Weighted distances**: supply weights (length p) for "minkowski" and "manhattan" (the latter uses p = 1). - Computations run in float32 and are JIT-compiled with JAX; where applicable, GPU/Metal/CPU device selection follows your existing backend.

Value

An $n \times m$ distance matrix in the format specified by return_type. If as_dist = TRUE and B = NULL (symmetric case), returns a [dist](#) object.

See Also

[dist](#)

Examples

```
## Not run:
# Simple Euclidean within-matrix distances (returns an n x n matrix)
X <- matrix(rnorm(50 * 8), 50, 8)
D <- fast_distance(X, metric = "euclidean")

# Cosine distance between two sets
A <- matrix(rnorm(100 * 16), 100, 16)
B <- matrix(rnorm(120 * 16), 120, 16)
Dcos <- fast_distance(A, B, metric = "cosine")

# Minkowski with p = 3 and feature weights
w <- runif(ncol(A))
Dm3 <- fast_distance(A, B, metric = "minkowski", p = 3, weights = w)

# Mahalanobis (diagonal approx, fast & robust)
Dmah_diag <- fast_distance(X, metric = "mahalanobis", approximate_inv = TRUE)

# Mahalanobis with full inverse (computed internally)
```

```

Dmah_full <- fast_distance(X, metric = "mahalanobis", approximate_inv = FALSE)

# Return a base R 'dist' object
D_dist <- fast_distance(X, metric = "euclidean", as_dist = TRUE)

## End(Not run)

```

find_fiducial_interval_R

Fiducial interval logic in base R, for randomization test

Description

Fiducial interval logic in base R, for randomization test

Usage

```

find_fiducial_interval_R(
  obsW,
  obsY,
  allW,
  tau_obs,
  alpha = 0.05,
  c_initial = 2,
  n_search_attempts = 500
)

```

Arguments

obsW	Observed assignment (0/1).
obsY	Observed outcome.
allW	Matrix of candidate random assignments (rows = assignments).
tau_obs	Observed difference in means with obsW, obsY.
alpha	Significance level (default 0.05).
c_initial	A numeric step scale (default 2).
n_search_attempts	Number of bracket search attempts (default 500).

Value

2-element numeric vector [lower, upper] or [NA, NA] if none accepted.

generate_randomizations

Generate randomizations for a rerandomization-based experimental design

Description

This function generates randomizations for experimental design using either exact enumeration or Monte Carlo sampling methods. It provides a unified interface to both approaches while handling memory and computational constraints appropriately.

Usage

```
generate_randomizations(
  n_units,
  n_treated,
  X = NULL,
  randomization_accept_prob,
  threshold_func = NULL,
  max_draws = 10^6,
  batch_size = 1000,
  randomization_type = "monte_carlo",
  approximate_inv = TRUE,
  file = NULL,
  return_type = "R",
  verbose = TRUE,
  conda_env = "fastrerandomize_env",
  conda_env_required = TRUE
)
```

Arguments

n_units	An integer specifying the total number of experimental units.
n_treated	An integer specifying the number of units to be assigned to treatment.
X	A numeric matrix of covariates used for balance checking. Cannot be NULL.
randomization_accept_prob	A numeric value between 0 and 1 specifying the probability threshold for accepting randomizations based on balance.
threshold_func	A 'JAX' function that computes a balance measure for each randomization. Only used for Monte Carlo sampling.
max_draws	An integer specifying the maximum number of randomizations to draw in Monte Carlo sampling.
batch_size	An integer specifying batch size for Monte Carlo processing.
randomization_type	A string specifying the type of randomization: either "exact" or "monte_carlo".
approximate_inv	A logical value indicating whether to use an approximate inverse (diagonal of the covariance matrix) instead of the full matrix inverse when computing balance metrics. This can speed up computations for high-dimensional covariates. Default is TRUE.

file	A string specifying where to save candidate randomizations (if saving, not returning).
return_type	A string specifying the format of the returned randomizations and balance measures. Allowed values are "R" for base R objects (e.g., matrix, numeric) or "jax" for 'JAX' arrays. Default is "R".
verbose	A logical value indicating whether to print progress information. Default is TRUE.
conda_env	A character string specifying the name of the conda environment to use via reticulate. Default is "fastrerandomize_env".
conda_env_required	A logical indicating whether the specified conda environment must be strictly used. If TRUE, an error is thrown if the environment is not found. Default is TRUE.

Details

The function supports two methods of generating randomizations:

1. Exact enumeration: Generates all possible randomizations (memory intensive but exact).
2. Monte Carlo sampling: Generates randomizations through sampling (more memory efficient).

For large problems (e.g., X with >20 rows), Monte Carlo sampling is recommended.

Value

Returns an S3 object with slots:

- assignments An array where each row represents one possible treatment assignment vector containing the accepted randomizations.
- balance_measures A numeric vector containing the balance measure for each corresponding randomization.
- fastrr_env The fastrerandomize environment.
- file_output If file is specified, results are saved to the given file path instead of being returned.

See Also

[generate_randomizations_exact](#) for the exact enumeration method. [generate_randomizations_mc](#) for the Monte Carlo sampling method.

Examples

```
## Not run:
# Generate synthetic data
X <- matrix(rnorm(20*5), 20, 5)

# Generate randomizations using exact enumeration
RandomizationSet_Exact <- generate_randomizations(
  n_units = nrow(X),
  n_treated = round(nrow(X)/2),
  X = X,
  randomization_accept_prob=0.1,
  randomization_type="exact")
```

```
# Generate randomizations using Monte Carlo sampling
RandomizationSet_MC <- generate_randomizations(
  n_units = nrow(X),
  n_treated = round(nrow(X)/2),
  X = X,
  randomization_accept_prob = 0.1,
  randomization_type = "monte_carlo",
  max_draws = 100000,
  batch_size = 1000)

## End(Not run)
```

```
generate_randomizations_exact
```

Generate Complete Randomizations with Optional Balance Constraints

Description

Generates all possible treatment assignments for a completely randomized experiment, optionally filtering them based on covariate balance criteria. The function can generate either all possible randomizations or a subset that meets specified balance thresholds using Hotelling's T-squared statistic.

Usage

```
generate_randomizations_exact(
  n_units,
  n_treated,
  X = NULL,
  randomization_accept_prob = 1,
  approximate_inv = TRUE,
  threshold_func = NULL,
  verbose = TRUE,
  conda_env = "fastrerandomize_env",
  conda_env_required = TRUE
)
```

Arguments

<code>n_units</code>	An integer specifying the total number of experimental units
<code>n_treated</code>	An integer specifying the number of units to be assigned to treatment
<code>X</code>	A numeric matrix of covariates where rows represent units and columns represent different covariates. Default is <code>NULL</code> , in which case all possible randomizations are returned without balance filtering.
<code>randomization_accept_prob</code>	A numeric value between 0 and 1 specifying the quantile threshold for accepting randomizations based on balance statistics. Default is 1 (accept all randomizations).

approximate_inv	A logical value indicating whether to use an approximate inverse (diagonal of the covariance matrix) instead of the full matrix inverse when computing balance metrics. This can speed up computations for high-dimensional covariates. Default is TRUE.
threshold_func	A function that calculates balance statistics for candidate randomizations. Default is VectorizedFastHotell12T2 which computes Hotelling's T-squared statistic.
verbose	A logical value indicating whether to print progress information. Default is TRUE.
conda_env	A character string specifying the name of the conda environment to use via reticulate. Default is "fastrerandomize_env".
conda_env_required	A logical indicating whether the specified conda environment must be strictly used. If TRUE, an error is thrown if the environment is not found. Default is TRUE.

Details

The function works in two main steps: 1. Generates all possible combinations of treatment assignments given `n_units` and `n_treated` 2. If covariates (`X`) are provided, filters these combinations based on balance criteria using the specified threshold function

The balance filtering process uses Hotelling's T-squared statistic by default to measure multivariate balance between treatment and control groups. Randomizations are accepted if their balance measure is below the specified quantile threshold.

Value

The function returns a *list* with two elements: `candidate_randomizations`: an array of randomization vectors `M_candidate_randomizations`: an array of their balance measures.

Note

This function requires 'JAX' and 'NumPy' to be installed and accessible through the reticulate package.

References

Hotelling, H. (1931). The generalization of Student's ratio. The Annals of Mathematical Statistics, 2(3), 360-378.

See Also

[generate_randomizations](#) for full randomization generation function. [generate_randomizations_mc](#) for the Monte Carlo version.

Examples

```
## Not run:
# Generate synthetic data
X <- matrix(rnorm(60), nrow = 10) # 10 units, 6 covariates

# Generate balanced randomizations with covariates
```

```
BalancedRandomizations <- generate_randomizations_exact(
  n_units = 10,
  n_treated = 5,
  X = X,
  randomization_accept_prob = 0.25 # Keep top 25% most balanced
)

## End(Not run)
```

```
generate_randomizations_mc
```

Draws a random sample of acceptable randomizations from all possible complete randomizations using Monte Carlo sampling

Description

This function performs sampling with replacement to generate randomizations in a memory-efficient way. It processes randomizations in batches to avoid memory issues and filters them based on covariate balance. The function uses JAX for fast computation and memory management.

Usage

```
generate_randomizations_mc(
  n_units,
  n_treated,
  X,
  randomization_accept_prob = 1,
  threshold_func = NULL,
  max_draws = 1e+05,
  batch_size = 1000,
  approximate_inv = TRUE,
  verbose = TRUE,
  conda_env = "fastrerandomize_env",
  conda_env_required = TRUE
)
```

Arguments

<code>n_units</code>	An integer specifying the total number of experimental units.
<code>n_treated</code>	An integer specifying the number of units to be assigned to treatment.
<code>X</code>	A numeric matrix of covariates used for balance checking. Cannot be NULL.
<code>randomization_accept_prob</code>	A numeric value between 0 and 1 specifying the probability threshold for accepting randomizations based on balance. Default is 1
<code>threshold_func</code>	A JAX function that computes a balance measure for each randomization. Must be vectorized using <code>jax.vmap</code> with <code>in_axes = list(NULL, 0L, NULL, NULL)</code> , and inputs covariates (matrix of X), treatment_assignment (vector of 0s and 1s), <code>n0</code> (scalar), <code>n1</code> (scalar). Default is <code>VectorizedFastHottel12T2</code> which uses Hotelling's T-squared statistic.

max_draws	An integer specifying the maximum number of randomizations to draw.
batch_size	An integer specifying how many randomizations to process at once. Lower values use less memory but may be slower.
approximate_inv	A logical value indicating whether to use an approximate inverse (diagonal of the covariance matrix) instead of the full matrix inverse when computing balance metrics. This can speed up computations for high-dimensional covariates. Default is TRUE.
verbose	A logical value indicating whether to print detailed information about batch processing progress, and GPU memory usage. Default is FALSE.
conda_env	A character string specifying the name of the conda environment to use via reticulate. Default is "fastrerandomize_env".
conda_env_required	A logical indicating whether the specified conda environment must be strictly used. If TRUE, an error is thrown if the environment is not found. Default is TRUE.

Details

The function works by:

1. Generating batches of random permutations.
2. Computing balance measures for each permutation using the provided threshold function.
3. Keeping only the top permutations that meet the acceptance probability threshold.
4. Managing memory by clearing unused objects and caches between batches.

The function uses smaller data types (int8, float16) where possible to reduce memory usage. It also includes assertions to verify array shapes and dimensions throughout.

Value

The function returns a *list* with two elements: `candidate_randomizations`: an array of randomization vectors `M_candidate_randomizations`: an array of their balance measures.

See Also

[generate_randomizations](#) for full randomization generation function. [generate_randomizations_exact](#) for the exact version.

Examples

```
## Not run:
# Generate synthetic data
X <- matrix(rnorm(100*5), 100, 5) # 5 covariates

# Generate 1000 randomizations for 100 units with 50 treated
rand_less_strict <- generate_randomizations_mc(
  n_units = 100,
  n_treated = 50,
  X = X,
  randomization_accept_prob=0.01,
  max_draws = 100000,
  batch_size = 1000)
```

```
# Use a stricter balance criterion
rand_more_strict <- generate_randomizations_mc(
  n_units = 100,
  n_treated = 50,
  X = X,
  randomization_accept_prob=0.001,
  max_draws = 1000000,
  batch_size = 1000)

## End(Not run)
```

```
generate_randomizations_R
```

Generate randomizations in base R, filtering by Hotelling's T^2 acceptance

Description

Base R function to either do exact enumeration or Monte Carlo random permutations, then keep the fraction whose T^2 is below the acceptance cutoff.

Usage

```
generate_randomizations_R(
  n_units,
  n_treated,
  X,
  accept_prob,
  random_type,
  max_draws,
  batch_size
)
```

Arguments

<code>n_units</code>	Integer, total number of units.
<code>n_treated</code>	Integer, number of units to be assigned to treatment.
<code>X</code>	Covariate matrix (<code>n_units</code> x <code>p</code>).
<code>accept_prob</code>	Numeric in [0, 1]: keep the fraction of randomizations that have the lowest T^2 up to this quantile.
<code>random_type</code>	Either "exact" or "monte_carlo".
<code>max_draws</code>	If 'random_type="monte_carlo"', how many permutations to sample.
<code>batch_size</code>	If 'random_type="monte_carlo"', how many permutations to handle per chunk.

Value

A list with:

- `randomizations`: a matrix (rows = accepted assignments).
- `balance`: numeric vector of T^2 values for each accepted assignment.

hotellingT2_R	<i>Compute Hotelling's T^2 in base R</i>
---------------	---

Description

This function provides a base R implementation of Hotelling's T^2 balance measure, renamed with `'_R'` for clarity that it is the R-based analog to the JAX version in `fastrerandomize`.

Usage

```
hotellingT2_R(X, W)
```

Arguments

X	A numeric matrix of covariates (observations in rows).
W	A 0/1 treatment assignment vector of the same length as <code>nrow(X)</code> .

Value

A numeric scalar: the Hotelling's T^2 for that assignment.

plot.fastrerandomize_randomizations	<i>Plot method for fastrerandomize_test objects</i>
-------------------------------------	---

Description

Plots the observed treatment effect and, if available, the fiducial interval on a horizontal axis.

Usage

```
## S3 method for class 'fastrerandomize_randomizations'
plot(x, ...)
```

Arguments

x	An object of class <code>fastrerandomize_test</code> .
...	Further graphical parameters passed to plot .

Value

No return value. This function is called for the side effect of generating a histogram of the accepted balance measures of object with class `fastrerandomize_randomizations`.

`plot.fastrerandomize_test`*Plot method for fastrerandomize_test objects*

Description

Plots a simple visualization of the observed effect and the fiducial interval (if present) on a horizontal axis.

Usage

```
## S3 method for class 'fastrerandomize_test'
plot(x, ...)
```

Arguments

<code>x</code>	An object of class <code>fastrerandomize_test</code> .
<code>...</code>	Further graphical parameters passed to plot .

Value

No output returned. Performs side effect of plotting `fastrerandomize_test` class objects.

`print.fastrerandomize_randomizations`*Print method for fastrerandomize_randomizations objects*

Description

Print method for `fastrerandomize_randomizations` objects

Usage

```
## S3 method for class 'fastrerandomize_randomizations'
print(x, ...)
```

Arguments

<code>x</code>	An object of class <code>fastrerandomize_instance</code> .
<code>...</code>	Further arguments passed to or from other methods.

Value

Prints an object of class `fastrerandomize_randomizations`.

```
print.fastrerandomize_test
```

Print method for fastrerandomize_test objects

Description

Print method for fastrerandomize_test objects

Usage

```
## S3 method for class 'fastrerandomize_test'
print(x, ...)
```

Arguments

x	An object of class fastrerandomize_test.
...	Further arguments passed to or from other methods.

Value

No return value, prints object of class fastrerandomize_test.

```
print2
```

Print timestamped messages with optional quieting

Description

This function prints messages prefixed with the current timestamp in a standardized format. Messages can be suppressed using the quiet parameter.

Usage

```
print2(text, quiet = FALSE)
```

Arguments

text	A character string containing the message to be printed.
quiet	A logical value indicating whether to suppress output. Default is FALSE.

Details

The function prepends the current timestamp in "YYYY-MM-DD HH:MM:SS" format to the provided message.

Value

No return value, called for side effect of printing with timestamp.

See Also

`Sys.time` for the underlying timestamp functionality.

Examples

```
# Print a basic message with timestamp
print2("Processing started")

# Suppress output
print2("This won't show", quiet = TRUE)

# Use in a loop
for(i in 1:3) {
  print2(sprintf("Processing item %d", i))
}
```

QJEData

QJEData: Agricultural Treatment Experiment Data

Description

Data from a field experiment studying moral hazard in tenancy contracts in agriculture.

After subsetting, this dataset includes observations on 968 experimental units with the following variables of interest: household composition, treatment assignment, and agricultural outcomes.

Usage

```
data(QJEData)
```

Format

A data frame with 968 rows and 7 columns:

children Numeric (integer). Number of children in the household. Larger numbers may reflect increased household labor needs and different investment or effort incentives.

married Numeric/binary. Whether the household head is currently married (1) or not (0). Marital status may influence decision-making and risk preferences in farming.

hh_size Numeric (integer). Household size. Differences in family labor availability or consumption needs can influence effort levels and thus relate to moral hazard in production decisions.

hh_sexrat Numeric. The ratio of adult men to adult women in the household. Imbalances in the male–female ratio can affect labor division and investment decisions.

treat1 Numeric/binary. Primary treatment indicator (e.g., whether a farmer is offered a specific tenancy contract or cost-sharing arrangement).

R_yield_ELA_sqm Numeric. Crop yield per square meter (e.g., kilograms of output per square meter). This is a principal outcome measure for evaluating productivity and treatment impact on farm performance.

ELA_Fertil_D Numeric/binary. Indicator for whether fertilizer was used (1) or not (0). This measures input investment—a key mechanism in moral hazard models (farmers may alter input use under different contracts).

Source

Burchardi, K.B., Ghatak, M., & Johanssen, A. (2019). Moral hazard: Experimental evidence from tenancy contracts. *The Quarterly Journal of Economics*, 134(1), 281-347.

randomization_test	<i>Fast randomization test</i>
--------------------	--------------------------------

Description

Fast randomization test

Usage

```
randomization_test(
  obsW = NULL,
  obsY = NULL,
  X = NULL,
  alpha = 0.05,
  candidate_randomizations = NULL,
  candidate_randomizations_array = NULL,
  n0_array = NULL,
  n1_array = NULL,
  randomization_accept_prob = 1,
  findFI = FALSE,
  c_initial = 2,
  max_draws = 10^6,
  batch_size = 10^5,
  randomization_type = "monte_carlo",
  approximate_inv = TRUE,
  file = NULL,
  verbose = TRUE,
  conda_env = "fastrerandomize_env",
  conda_env_required = TRUE
)
```

Arguments

obsW	A numeric vector where 0's correspond to control units and 1's to treated units.
obsY	An optional numeric vector of observed outcomes. If not provided, the function assumes a NULL value.
X	A numeric matrix of covariates.
alpha	The significance level for the test. Default is 0.05.
candidate_randomizations	A numeric matrix of candidate randomizations.
candidate_randomizations_array	An optional 'JAX' array of candidate randomizations. If not provided, the function coerces candidate_randomizations into a 'JAX' array.
n0_array	An optional array specifying the number of control units.
n1_array	An optional array specifying the number of treated units.

randomization_accept_prob	An numeric scalar or vector of probabilities for accepting each randomization.
findFI	A logical value indicating whether to find the fiducial interval. Default is FALSE.
c_initial	A numeric value representing the initial criterion for the randomization. Default is 2.
max_draws	An integer specifying the maximum number of candidate randomizations to generate (or to consider) for the test when randomization_type = "monte_carlo". Default is 1e6.
batch_size	An integer specifying the batch size for Monte Carlo sampling. Batches are processed one at a time for memory efficiency. Default is 1e5.
randomization_type	A string specifying the type of randomization for the test. Allowed values are "exact" or "monte_carlo". Default is "monte_carlo".
approximate_inv	A logical value indicating whether to use an approximate inverse (diagonal of the covariance matrix) instead of the full matrix inverse when computing balance metrics. This can speed up computations for high-dimensional covariates. Default is TRUE.
file	A character string specifying the path (including filename) where candidate randomizations will be saved or loaded from. If NULL, randomizations remain in memory. Default is NULL.
verbose	A logical value indicating whether to print progress information. Default is TRUE.
conda_env	A character string specifying the name of the conda environment to use via reticulate. Default is "fastrerandomize_env".
conda_env_required	A logical indicating whether the specified conda environment must be strictly used. If TRUE, an error is thrown if the environment is not found. Default is TRUE.

Value

Returns an S3 object with slots:

- `p_value` A numeric value or vector representing the p-value of the test (or the expected p-value under the prior structure specified in the function inputs).
- `FI` A numeric vector representing the fiducial interval if `findFI=TRUE`.
- `tau_obs` A numeric value or vector representing the estimated treatment effect(s).
- `fastrr_env` The `fastrerandomize` environment.

References

- Zhang, Y. and Zhao, Q., 2023. What is a randomization test?. Journal of the American Statistical Association, 118(544), pp.2928-2942.

See Also

[generate_randomizations](#) for randomization generation function.

Examples

```
## Not run:
# A small synthetic demonstration with 6 units, 3 treated and 3 controls:

# Generate pre-treatment covariates
X <- matrix(rnorm(24*2), ncol = 2)

# Generate candidate randomizations
RandomizationSet_MC <- generate_randomizations(
  n_units = nrow(X),
  n_treated = round(nrow(X)/2),
  X = X,
  randomization_accept_prob = 0.1,
  randomization_type = "monte_carlo",
  max_draws = 100000,
  batch_size = 1000
)

# Generate outcome
W <- RandomizationSet_MC$randomizations[1,]
obsY <- rnorm(nrow(X), mean = 2 * W)

# Perform randomization test
results_base <- randomization_test(
  obsW = W,
  obsY = obsY,
  X = X,
  candidate_randomizations = RandomizationSet_MC$randomizations,
)
print(results_base)

# Perform randomization test
result_fi <- randomization_test(
  obsW = W,
  obsY = obsY,
  X = X,
  candidate_randomizations = RandomizationSet_MC$randomizations,
  findFI = TRUE
)
print(result_fi)

## End(Not run)
```

randomization_test_R	<i>Base R randomization test: difference in means + optional fiducial interval</i>
----------------------	--

Description

Base R randomization test: difference in means + optional fiducial interval

Usage

```
randomization_test_R(obsW, obsY, allW, findFI = FALSE, alpha = 0.05)
```

Arguments

obsW	Observed assignment (0/1).
obsY	Observed outcome vector.
allW	Matrix of candidate random assignments (rows = assignments).
findFI	Logical, whether to compute fiducial interval as well.
alpha	Significance level (default 0.05).

Value

A list with p_value, tau_obs, and (optionally) FI if 'findFI=TRUE'.

```
summary.fastrerandomize_randomizations
```

Summary method for fastrerandomize_randomizations objects

Description

Summary method for fastrerandomize_randomizations objects

Usage

```
## S3 method for class 'fastrerandomize_randomizations'
summary(object, ...)
```

Arguments

object	An object of class fastrerandomize_randomizations.
...	Further arguments passed to or from other methods.

Value

A list with summary statistics, printed by default.

```
summary.fastrerandomize_test
```

Summary method for fastrerandomize_test objects

Description

Summary method for fastrerandomize_test objects

Usage

```
## S3 method for class 'fastrerandomize_test'
summary(object, ...)
```

Arguments

- object An object of class `fastrerandomize_test`.
- ... Further arguments passed to or from other methods.

Value

Returns an (invisible) list with a summary of `fastrerandomize_test` class objects.

YOPData	<i>YOPData</i>
---------	----------------

Description

Data from a re-analysis of the Youth Opportunities Program anti-poverty RCT in Uganda, with satellite imagery neural representations linked to RCT units.

Usage

`data(YOPData)`

Format

A list containing two data frames:

- RCTData** Treatment, outcome, and geolocation information
- ImageEmbeddings** CLIP-RSICD neural embeddings of satellite imagery

Source

- Blattman, C., Fiala, N. and Martinez, S. (2020). The Long-term Impacts of Grants on Poverty: Nine-year Evidence from Uganda’s Youth Opportunities Program. *American Economic Review: Insights*, 2(3), 287-304.
- Jerzak, C.T., Johansson, F.D. and Daoud, A. (2023). Image-based Treatment Effect Heterogeneity. *Conference on Causal Learning and Reasoning*, 531-552. PMLR.

Index

build_backend, [2](#)

check_jax_availability, [3](#)
compute_diff_at_tau_for_oneW_R, [4](#)

diff_in_means_R, [4](#)
dist, [7](#)

fast_distance, [6](#)
fastrerandomize_class, [5](#)
fastrerandomize_test, [5](#)
find_fiducial_interval_R, [8](#)

generate_randomizations, [5](#), [9](#), [12](#), [14](#), [21](#)
generate_randomizations_exact, [10](#), [11](#),
[14](#)
generate_randomizations_mc, [10](#), [12](#), [13](#)
generate_randomizations_R, [15](#)

hotellingT2_R, [16](#)

plot, [16](#), [17](#)
plot.fastrerandomize_randomizations,
[16](#)
plot.fastrerandomize_test, [17](#)
print.fastrerandomize_randomizations,
[17](#)
print.fastrerandomize_test, [18](#)
print2, [18](#)

QJEData, [19](#)

randomization_test, [20](#)
randomization_test_R, [22](#)

summary.fastrerandomize_randomizations,
[23](#)
summary.fastrerandomize_test, [23](#)

YOPData, [24](#)